

Sentiment_Classification_Comparison

June 13, 2020

1 Sentiment binary classification through just one hidden layer definition using Python

2 Summary of binary sentiment classification process:

- Create three Counter objects to store positive, negative and total counts
- Normalize for the effect of common words like the, is, ... that do not convey any sentiment load through calculating `np.log(pos_to_neg_ratio)` counter to focus on the words found in positive reviews more often than in negative reviews, and vice versa. As a result, neutral words will be close to 0, words will get more positive as their ratios approach and go above 1, and words will get more negative as their ratios approach and go below -1.
- Tokenize review and label words by transforming them into numbers: >* Create a set of unique words from all reviews as the vocabulary set (use set) >* Map each unique words to an integer (use enumerate) >* Map labels to 0 or 1.
- Reduce noise by strategically reducing the Vocabulary such that important sentiments stand out: >* Apply a `min_count` so that only words higher than cut-off are added to the vocabulary (avoid rare words and reduce noise). >* Apply `apolarity_cutoff` so that only words with their postive-to-negative higher than the cut-off are added to the vocabulary (discard neutral words as possible).
- Build the network efficiently: >* Find the indecis of all words in the input review from the processed training vocabulary. >* Update hidden layer just by updating the weights of elements that correspond to the indecis of the input words (discrad tonnes of unnecessary matrix multiplication)

```
[2]: def pretty_print_review_and_label(i):
      print(labels[i] + "\t:\t" + reviews[i][:80] + "...") # [:80] limits the
      ↪number of characters of each review per line

g = open('data/reviews.txt','r') # What we know!
reviews = list(map(lambda x:x[:-1].lower(),g.readlines())) # .lower() methods
      ↪makes all words lower case for consistent treatment.
g.close()

g = open('data/labels.txt','r') # What we WANT to know!
labels = list(map(lambda x:x[:-1].upper(),g.readlines()))
```

```
g.close()
```

```
[3]: import time
import sys
from collections import Counter # Counter is a convenient fast dictionary that
    ↪ already includes the original keys
import numpy as np

# Encapsulate our neural network in a class
class SentimentNetwork:
    def __init__(self, reviews, labels, min_count = 10, polarity_cutoff = 0.1,
    ↪ hidden_nodes = 10, learning_rate = 0.1):
        """Create a SentimenNetwork with the given settings
        Args:
            reviews(list) - List of reviews used for training
            labels(list) - List of POSITIVE/NEGATIVE labels associated with the
    ↪ given reviews
            min_count(int) - Words should only be added to the vocabulary
                           if they occur more than this many times
            polarity_cutoff(float) - The absolute value of a word's
    ↪ positive-to-negative
                                   ratio must be at least this big to be
    ↪ considered.
            hidden_nodes(int) - Number of nodes to create in the hidden layer
            learning_rate(float) - Learning rate to use while training

        """
        # Assign a seed to our random number generator to ensure we get
        # reproducible results during development
        np.random.seed(1)

        # process the reviews and their associated labels so that everything
        # is ready for training
        self.pre_process_data(reviews, labels, polarity_cutoff, min_count)

        # Build the network to have the number of hidden nodes and the learning
    ↪ rate that
        # were passed into this initializer. Make the same number of input
    ↪ nodes as
        # there are vocabulary words and create a single output node.
        self.init_network(len(self.review_vocab), hidden_nodes, 1, learning_rate)

    def pre_process_data(self, reviews, labels, polarity_cutoff, min_count):
        ## Calculate positive-to-negative ratios for words before building
    ↪ vocabulary
```

```

positive_counts = Counter()
negative_counts = Counter()
total_counts = Counter()

for i in range(len(reviews)):
    if(labels[i] == 'POSITIVE'):
        for word in reviews[i].split(" "):
            positive_counts[word] += 1
            total_counts[word] += 1
    else:
        for word in reviews[i].split(" "):
            negative_counts[word] += 1
            total_counts[word] += 1

pos_neg_ratios = Counter()

for term,cnt in list(total_counts.most_common()):
    if(cnt >= 50): # the pos_neg_ratio ratio is not calculatd for a
    ↪term that happens less than this count in a huge dataset
                    # (regarless of being actually very positive or
    ↪negative in few reviwes) .
        pos_neg_ratio = positive_counts[term] /
    ↪float(negative_counts[term]+1)
        pos_neg_ratios[term] = pos_neg_ratio

for word,ratio in pos_neg_ratios.most_common():
    if(ratio > 1):
        pos_neg_ratios[word] = np.log(ratio)
    else:
        pos_neg_ratios[word] = -np.log((1 / (ratio + 0.01)))

# populate review_vocab with all of the words in the given reviews.
review_vocab = set()
for review in reviews:
    for word in review.split(" "):
        # only add words that occur at least min_count times
        # and for words with pos/neg ratios, only add words that meet
    ↪the polarity_cutoff
        if(total_counts[word] > min_count):
            if(word in pos_neg_ratios.keys()): # recall that for a word
    ↪to be in these pos_neg_ratios.keys(),
                                                    # it should have
    ↪repeated more than a specific cnt!
                if((pos_neg_ratios[word] >= polarity_cutoff) or
    ↪(pos_neg_ratios[word] <= -polarity_cutoff)):
                    review_vocab.add(word)

```

```

        else:
            review_vocab.add(word)

    # Convert the vocabulary set to a list so we can access words via
    ↪ indices
    self.review_vocab = list(review_vocab)

    # populate label_vocab with all of the words in the given labels.
    label_vocab = set()
    for label in labels:
        label_vocab.add(label)

    # Convert the label vocabulary set to a list so we can access labels
    ↪ via indices
    self.label_vocab = list(label_vocab)

    # Store the sizes of the review and label vocabularies.
    self.review_vocab_size = len(self.review_vocab)
    self.label_vocab_size = len(self.label_vocab)

    # Create a dictionary of words in the vocabulary mapped to index
    ↪ positions
    self.word2index = {}
    for index, word in enumerate(self.review_vocab):
        self.word2index[word] = index # populate self.word2index with
    ↪ indices for all the words in self.review_vocab

    # Create a dictionary of labels mapped to index positions
    self.label2index = {}
    for index, label in enumerate(self.label_vocab):
        self.label2index[label] = index

    def init_network(self, input_nodes, hidden_nodes, output_nodes,
    ↪ learning_rate):
        # Store the number of nodes in input, hidden, and output layers.
        self.input_nodes = input_nodes
        self.hidden_nodes = hidden_nodes
        self.output_nodes = output_nodes

        # Store the learning rate
        self.learning_rate = learning_rate

        # Initialize weights

        # initialize self.weights_0_1 as a matrix of zeros.
        # These are the weights between the input layer and the hidden layer.
        self.weights_0_1 = np.zeros(shape=(self.input_nodes, self.hidden_nodes))

```

```

        # initialize self.weights_1_2 as a matrix of random values.
        # These are the weights between the hidden layer and the output layer.
        self.weights_1_2 = np.random.normal(loc=0.0, scale=self.
→hidden_nodes**-0.5, size=(self.hidden_nodes, self.output_nodes))

        # The input layer, a two-dimensional matrix with shape 1 x hidden_nodes.
        self.layer_1 = np.zeros((1, hidden_nodes))

def get_target_for_label(self, label):
    if (label == 'POSITIVE'):
        return 1
    else:
        return 0

def sigmoid(self, x):
    return 1 / (1 + np.exp(-x))

def sigmoid_output_2_derivative(self, output):
    # Return the derivative of the sigmoid activation function,
    # where "output" is the original output from the sigmoid function
    return output * (1 - output)

def train(self, training_reviews_raw, training_labels):
    ## pre-process training reviews so we can deal directly with the
→indices of non-zero inputs
    training_reviews = list()
    for review in training_reviews_raw:
        indices = set()
        for word in review.split(" "):
            # if the particular word in the review exists in word2index
→vocabulary,
            # find its index and add the word index to the indices set
            if (word in self.word2index.keys()):
                indices.add(self.word2index[word]) # add method works with
→set
            # at the end of the loop, indices of all words are found from
→word2index vocabulary.
            # indices set is converted to list and its elements are appended to
→training_reviews list.
            # TAKE IMPORTANT NOTE that training_reviews is just a list of lists
→where each list is indices of words for a review.
            training_reviews.append(list(indices))

    # make sure out we have a matching number of reviews and labels
    assert(len(training_reviews) == len(training_labels))

```

```

# Keep track of correct predictions to display accuracy during training
correct_so_far = 0

# Remember when we started for printing time statistics
start = time.time()

# loop through all the given reviews and run a forward and backward
→pass,
# updating weights for every item
for i in range(len(training_reviews)):

    # Get the next review and its correct label
    review = training_reviews[i]
    label = labels[i]

    ##### Implement the forward pass here #####
    ### Forward pass ###

    # Hidden layer
    # no activation function to preserve the linearity
    ## Add in only the weights for non-zero items
    self.layer_1 *= 0
    for index in review:
        self.layer_1 += self.weights_0_1[index]

    # Output layer
    layer_2 = self.sigmoid(np.matmul(self.layer_1, self.weights_1_2))

    ##### Implement the backward pass here #####
    ### Backward pass ###

    # Output error
    layer_2_error = self.get_target_for_label(label) - layer_2
    layer_2_delta = layer_2_error * self.
→sigmoid_output_2_derivative(layer_2) # adjust for the slope of non-linearity

    # Backpropagated error
    layer_1_error = np.matmul(layer_2_delta, self.weights_1_2.T) #
→errors propagated to the hidden layer
    # note in feedforward weights_1_2 is multiplied with layer_1 as
→input while in backpropagation its transpose
    # is multiplied with downstream (next layer) error term as the
→input.
    layer_1_delta = layer_1_error * 1 # hidden layer gradients - no
→nonlinearity so it's the same as the error (no adjustment)

```

```

        # update weights
        self.weights_1_2 += self.learning_rate * layer_2_delta * self.
→ layer_1.T # update hidden-to-output weights with gradient descent step

        ## Only update the weights that were used in the forward pass
        for index in review:
            self.weights_0_1[index] += self.learning_rate *
→ layer_1_delta[0] # update input-to-hidden weights with gradient descent step

        # Keep track of correct label predictions
        # how accurate are the predictions
        if abs(layer_2_error) < 0.5:
            correct_so_far += 1

        # how fast we are training
        elapsed_time = float(time.time() - start)
        reviews_per_second = i / elapsed_time if elapsed_time > 0 else 0

        sys.stdout.write("\rProgress:" + str(100 * i /
→ float(len(training_reviews)))[:4] \
                        + "% Speed(reviews/sec):" +
→ str(reviews_per_second)[0:5] \
                        + " #Correct:" + str(correct_so_far) + " #Trained:
→ " + str(i+1) \
                        + " Training Accuracy:" + str(correct_so_far * 100
→ / float(i+1))[:4] + "%")
        if(i % 2500 == 0): # every 2500, make a new line
            print("")

    def test(self, testing_reviews, testing_labels):
        """
        Attempts to predict the labels for the given testing_reviews,
        and uses the test_labels to calculate the accuracy of those predictions.
        """

        # keep track of how many correct predictions we make
        correct = 0

        # we'll time how many predictions per second we make
        start = time.time()

        # Loop through each of the given reviews and call run to predict its
→ label.
        for i in range(len(testing_reviews)):
            pred = self.run(testing_reviews[i])

```

```

        if(pred == testing_labels[i]):
            correct += 1

        # For debug purposes, print out our prediction accuracy and speed
        # throughout the prediction process.

        elapsed_time = float(time.time() - start)
        reviews_per_second = i / elapsed_time if elapsed_time > 0 else 0

        sys.stdout.write("\rProgress:" + str(100 * i/
↪float(len(testing_reviews)))[:4] \
                                + "% Speed(reviews/sec):" + _
↪str(reviews_per_second)[0:5] \
                                + " #Correct:" + str(correct) + " #Tested:" + _
↪str(i+1) \
                                + " Testing Accuracy:" + str(correct * 100 / _
↪float(i+1))[:4] + "%")

    def run(self, review):
        """
        Returns a POSITIVE or NEGATIVE prediction for the given review.
        """

        # Run a forward pass through the network, like in the "train" function.

        # Hidden Layer
        ## Identify the indices used in the review and then add just those
↪weights to layer_1
        self.layer_1 *= 0
        unique_indices = set()
        for word in review.lower().split(" "):
            if word in self.word2index.keys():
                unique_indices.add(self.word2index[word])
        for index in unique_indices:
            self.layer_1 += self.weights_0_1[index]

        # Output Layer
        layer_2 = self.sigmoid(np.matmul(self.layer_1, self.weights_1_2))

        # Return POSITIVE for values above greater-than-or-equal-to 0.5 in the
↪output layer;
        # return NEGATIVE for other values

        if(layer_2[0] >= 0.5):
            return "POSITIVE"
        else:

```



```
return "NEGATIVE"
```

```
[4]: split_frac = 0.8

## split data into training, validation, and test data (features and labels, x
↪and y)

split_idx = int(len(reviews)*split_frac)
train_x, remaining_x = reviews[:split_idx], reviews[split_idx:]
train_y, remaining_y = labels[:split_idx], labels[split_idx:]

test_idx = int(len(remaining_x)*0.5)
val_x, test_x = remaining_x[:test_idx], remaining_x[test_idx:]
val_y, test_y = remaining_y[:test_idx], remaining_y[test_idx:]
```

```
[12]: # instantiate the model
mlp_full = SentimentNetwork(train_x,train_y,min_count=0,polarity_cutoff=0.
↪0,learning_rate=0.01)

# start the clock
start_time = time.time()
# call the etrain function

mlp_full.train(train_x, train_y)
# print the training time
print("\nTotal training time is %s minutes." % round((time.time() - start_time),
↪/ 60, 3))
```

```
Progress:0.0% Speed(reviews/sec):0 #Correct:0 #Trained:1 Training Accuracy:0.0%
Progress:12.5% Speed(reviews/sec):1439. #Correct:1943 #Trained:2501 Training
Accuracy:77.6%
Progress:25.0% Speed(reviews/sec):1386. #Correct:4002 #Trained:5001 Training
Accuracy:80.0%
Progress:37.5% Speed(reviews/sec):1390. #Correct:6118 #Trained:7501 Training
Accuracy:81.5%
Progress:50.0% Speed(reviews/sec):1398. #Correct:8272 #Trained:10001 Training
Accuracy:82.7%
Progress:62.5% Speed(reviews/sec):1399. #Correct:10432 #Trained:12501 Training
Accuracy:83.4%
Progress:75.0% Speed(reviews/sec):1398. #Correct:12553 #Trained:15001 Training
Accuracy:83.6%
Progress:87.5% Speed(reviews/sec):1397. #Correct:14677 #Trained:17501 Training
Accuracy:83.8%
Progress:99.9% Speed(reviews/sec):1390. #Correct:16867 #Trained:20000 Training
Accuracy:84.3%
Total training time is 0.28 minutes.
```

```
[14]: mlp_full.test(val_x, val_y)
```

Progress:99.9% Speed(reviews/sec):2287. #Correct:2194 #Tested:2500 Testing Accuracy:87.7%

```
[15]: mlp_full.test(test_x, test_y)
```

Progress:99.9% Speed(reviews/sec):2309. #Correct:2158 #Tested:2500 Testing Accuracy:86.3%

3 Alternative solution: Sentiment Analysis with an RNN - LSTM

```
[160]: # read data from text files
with open('data/reviews.txt', 'r') as f:
    reviews = f.read()
with open('data/labels.txt', 'r') as f:
    labels = f.read()
```

3.0.1 Data pre-processing

```
[161]: from string import punctuation
from collections import Counter

# get rid of punctuation
reviews = reviews.lower() # lowercase, standardize
all_text = ''.join([c for c in reviews if c not in punctuation])

# split by new lines and spaces
reviews_split = all_text.split('\n')
all_text = ' '.join(reviews_split)

# create a list of words
words = all_text.split()

## Encoding the words
## Build a dictionary that maps words to integers
counts = Counter(words)
vocab = sorted(counts, key=counts.get, reverse=True) # the most common word
↳ mapped to the integer value 1
vocab_to_int = {word: ii for ii, word in enumerate(vocab, 1)} # start from 1
↳ explicitly stated

## use the dict to tokenize each review in reviews_split
## store the tokenized reviews in reviews_ints
reviews_ints = []
for review in reviews_split:
```

```

    reviews_ints.append([vocab_to_int[word] for word in review.split()])

## Encoding the labels
# 1=positive, 0=negative label conversion
labels_split = labels.split('\n')
encoded_labels = np.array([1 if label == 'positive' else 0 for label in_
    ↪labels_split])

## Removing Outliers
## remove any reviews/labels with zero length from the reviews_ints list.

# get indices of any reviews with length of not 0 that we want to keep
non_zero_idx = [ii for ii, review in enumerate(reviews_ints) if len(review) !=_
    ↪0]

# remove 0-length reviews and their labels
reviews_ints = [reviews_ints[ii] for ii in non_zero_idx]
encoded_labels = np.array([encoded_labels[ii] for ii in non_zero_idx])

## Padding sequences
def pad_features(reviews_ints, seq_length):
    ''' Return features of review_ints, where each review is padded with 0's
    or truncated to the input seq_length.
    '''

    # getting the correct rows x cols shape
    features = np.zeros((len(reviews_ints), seq_length), dtype=int)

    # for each review, I grab that review and
    for i, row in enumerate(reviews_ints):
        features[i, -len(row):] = np.array(row)[:seq_length]

    return features

```

3.0.2 Training, Validation, Test

```

[162]: split_frac = 0.8

## split data into training, validation, and test data (features and labels, x_
    ↪and y)

split_idx = int(len(features) * split_frac)
train_x, remaining_x = features[:split_idx], features[split_idx:]
train_y, remaining_y = encoded_labels[:split_idx], encoded_labels[split_idx:]

valid_frac = 0.5

```

```
test_idx = int(len(remaining_x) * valid_frac)
val_x, test_x = remaining_x[:test_idx], remaining_x[test_idx:]
val_y, test_y = remaining_y[:test_idx], remaining_y[test_idx:]
```

```
[163]: ## print out the shapes of your resultant feature data
print("\t\t\tFeature Shapes:")
print("Train set: \t\t{}".format(train_x.shape),
      "\nValidation set: \t{}".format(val_x.shape),
      "\nTest set: \t\t{}".format(test_x.shape))
```

```

                                Feature Shapes:
Train set:                      (20000, 200)
Validation set:                  (2500, 200)
Test set:                       (2500, 200)
```

3.0.3 DataLoaders and Batching

```
[164]: import torch
from torch.utils.data import TensorDataset, DataLoader

# create Tensor datasets
train_data = TensorDataset(torch.from_numpy(train_x), torch.from_numpy(train_y))
valid_data = TensorDataset(torch.from_numpy(val_x), torch.from_numpy(val_y))
test_data = TensorDataset(torch.from_numpy(test_x), torch.from_numpy(test_y))

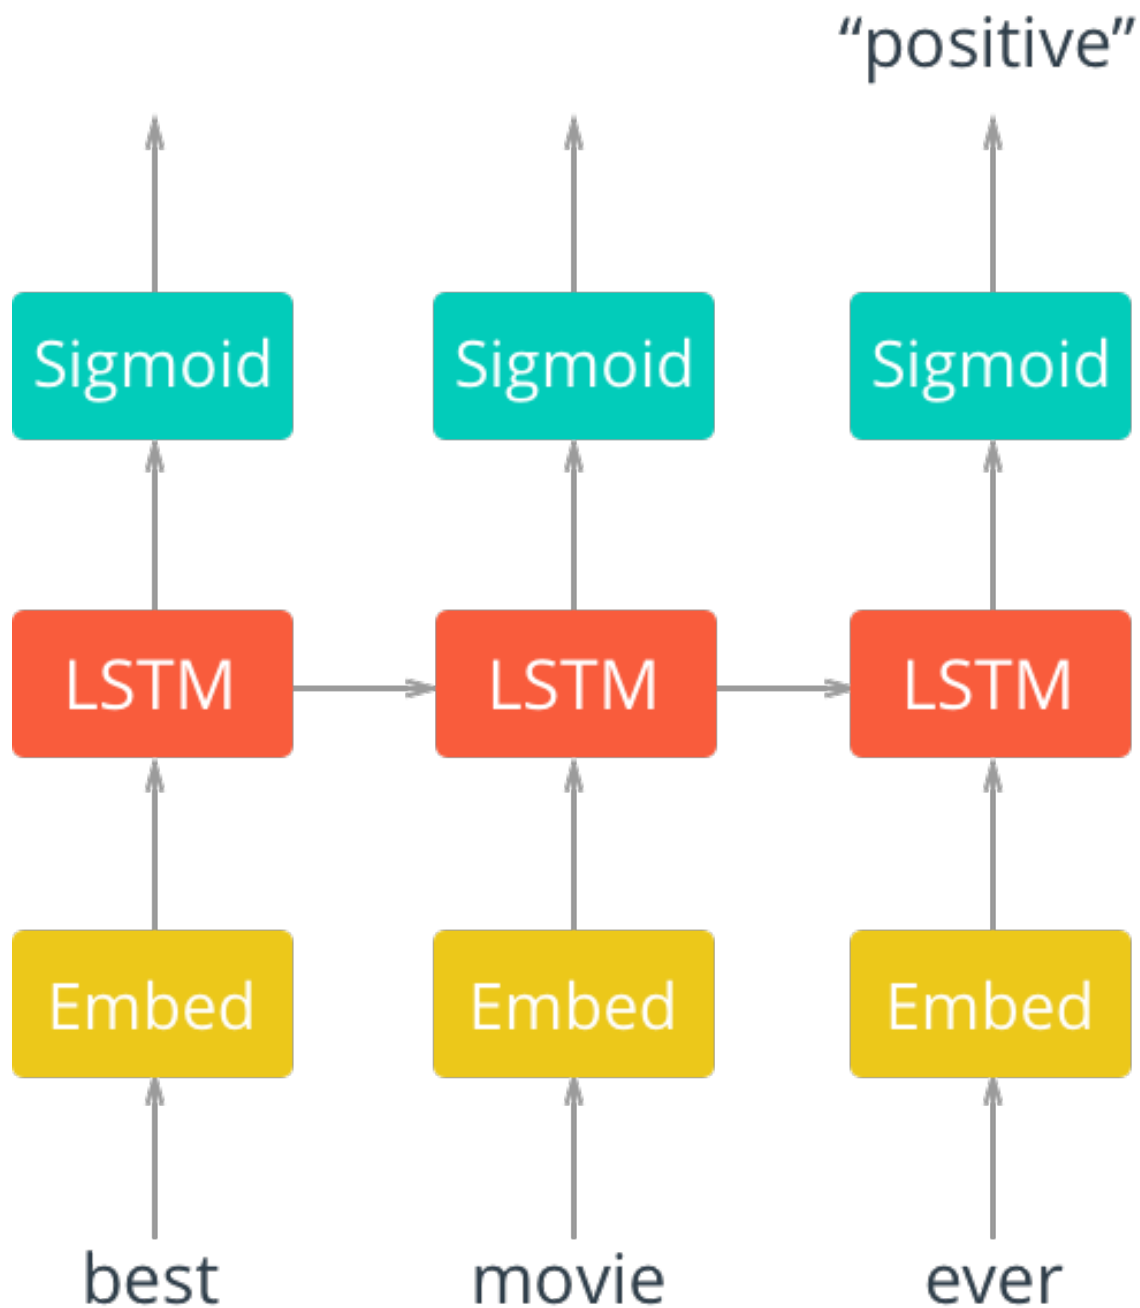
# dataloaders
batch_size = 50

# make sure to SHUFFLE your data
train_loader = DataLoader(train_data, shuffle=True, batch_size=batch_size)
valid_loader = DataLoader(valid_data, shuffle=True, batch_size=batch_size)
test_loader = DataLoader(test_data, shuffle=True, batch_size=batch_size)
```

4 Sentiment Network with PyTorch

```
[130]: from IPython.display import Image
Image(filename='assets/network_diagram.png')
```

[130]:



```
[131]: # First checking if GPU is available
train_on_gpu=torch.cuda.is_available()

if(train_on_gpu):
    print('Training on GPU.')
else:
    print('No GPU available, training on CPU.')
```

No GPU available, training on CPU.

4.0.1 Model definition

```
[165]: import torch.nn as nn

class SentimentRNN(nn.Module):
    """
    The RNN model that will be used to perform Sentiment analysis.
    """

    def __init__(self, vocab_size, output_size, embedding_dim, hidden_dim, n_layers, drop_prob=0.5):
        """
        Initialize the model by setting up the layers.
        """
        super(SentimentRNN, self).__init__()

        self.output_size = output_size
        self.n_layers = n_layers
        self.hidden_dim = hidden_dim

        # embedding and LSTM layers
        self.embedding = nn.Embedding(vocab_size, embedding_dim)
        self.lstm = nn.LSTM(embedding_dim, hidden_dim, n_layers,
                             dropout=drop_prob, batch_first=True)

        # dropout layer
        self.dropout = nn.Dropout(0.3)

        # linear and sigmoid layers
        self.fc = nn.Linear(hidden_dim, output_size)
        self.sig = nn.Sigmoid()

    def forward(self, x, hidden):
        """
        Perform a forward pass of our model on some input and hidden state.
        """
        batch_size = x.size(0)

        # embeddings and lstm_out
        x = x.long()
        embeds = self.embedding(x)
        lstm_out, hidden = self.lstm(embeds, hidden)

        # stack up lstm outputs
```

```

lstm_out = lstm_out.contiguous().view(-1, self.hidden_dim)

# dropout and fully-connected layer
out = self.dropout(lstm_out)
out = self.fc(out)
# sigmoid function
sig_out = self.sig(out)

# reshape to be batch_size first
sig_out = sig_out.view(batch_size, -1)
sig_out = sig_out[:, -1] # get last batch of labels

# return last sigmoid output and hidden state
return sig_out, hidden

def init_hidden(self, batch_size):
    ''' Initializes hidden state '''
    # Create two new tensors with sizes n_layers x batch_size x hidden_dim,
    # initialized to zero, for hidden state and cell state of LSTM
    weight = next(self.parameters()).data

    if (train_on_gpu):
        hidden = (weight.new(self.n_layers, batch_size, self.hidden_dim).
→zero_().cuda(),
                  weight.new(self.n_layers, batch_size, self.hidden_dim).
→zero_().cuda())
    else:
        hidden = (weight.new(self.n_layers, batch_size, self.hidden_dim).
→zero_(),
                  weight.new(self.n_layers, batch_size, self.hidden_dim).
→zero_())

    return hidden

```

```

[166]: # Instantiate the model w/ hyperparams
vocab_size = len(vocab_to_int)+1 # our word tokens +1 for the 0-token added to
→pad input features +
output_size = 1 # a single sigmoid value between 0 and 1
embedding_dim = 400
hidden_dim = 256
n_layers = 2

net = SentimentRNN(vocab_size, output_size, embedding_dim, hidden_dim, n_layers)

print(net)

```

```

SentimentRNN(
    (embedding): Embedding(74073, 400)
    (lstm): LSTM(400, 256, num_layers=2, batch_first=True, dropout=0.5)
    (dropout): Dropout(p=0.3, inplace=False)
    (fc): Linear(in_features=256, out_features=1, bias=True)
    (sig): Sigmoid()
)

```

4.0.2 Training

```

[167]: # loss and optimization functions
lr=0.001

```

```

criterion = nn.BCELoss()
optimizer = torch.optim.Adam(net.parameters(), lr=lr)

```

```

[169]: # training params
start = time.time()
epochs = 4 # 3-4 is approx where I noticed the validation loss stop decreasing

counter = 0
print_every = 100
clip=5 # gradient clipping

# move model to GPU, if available
if(train_on_gpu):
    net.cuda()

net.train()
# train for some number of epochs
for e in range(epochs):
    # initialize hidden state
    h = net.init_hidden(batch_size)

    # batch loop
    for inputs, labels in train_loader:
        counter += 1

        if(train_on_gpu):
            inputs, labels = inputs.cuda(), labels.cuda()

        # Creating new variables for the hidden state, otherwise
        # we'd backprop through the entire training history
        h = tuple([each.detach() for each in h])

        # zero accumulated gradients
        net.zero_grad()

```



```

# get the output from the model
output, h = net(inputs, h)

# calculate the loss and perform backprop
loss = criterion(output.squeeze(), labels.float())
loss.backward()
# `clip_grad_norm` helps prevent the exploding gradient problem in RNNs
↪ / LSTMs.
nn.utils.clip_grad_norm_(net.parameters(), clip)
optimizer.step()

# loss stats
if counter % print_every == 0:
    # Get validation loss
    val_h = net.init_hidden(batch_size)
    val_losses = []
    net.eval()
    for inputs, labels in valid_loader:

        # Creating new variables for the hidden state, otherwise
        # we'd backprop through the entire training history
        val_h = tuple([each.detach() for each in val_h])

        if(train_on_gpu):
            inputs, labels = inputs.cuda(), labels.cuda()

        output, val_h = net(inputs, val_h)
        val_loss = criterion(output.squeeze(), labels.float())

        val_losses.append(val_loss.item())

    net.train()
    print("Epoch: {}/{}...".format(e+1, epochs),
          "Step: {}...".format(counter),
          "Loss: {:.6f}...".format(loss.item()),
          "Val Loss: {:.6f}".format(np.mean(val_losses)))

print("\nTotal training time is %s minutes." % round((time.time() - start_time),
↪ / 60, 3))

```

```

Epoch: 1/4... Step: 100... Loss: 0.574357... Val Loss: 0.592211
Epoch: 1/4... Step: 200... Loss: 0.594059... Val Loss: 0.663425
Epoch: 1/4... Step: 300... Loss: 0.523350... Val Loss: 0.585923
Epoch: 1/4... Step: 400... Loss: 0.554656... Val Loss: 0.571660
Epoch: 2/4... Step: 500... Loss: 0.521586... Val Loss: 0.575272
Epoch: 2/4... Step: 600... Loss: 0.357778... Val Loss: 0.509176

```

```
Epoch: 2/4... Step: 700... Loss: 0.385232... Val Loss: 0.525697
Epoch: 2/4... Step: 800... Loss: 0.406087... Val Loss: 0.461405
Epoch: 3/4... Step: 900... Loss: 0.305346... Val Loss: 0.490994
Epoch: 3/4... Step: 1000... Loss: 0.224487... Val Loss: 0.455219
Epoch: 3/4... Step: 1100... Loss: 0.340713... Val Loss: 0.462203
Epoch: 3/4... Step: 1200... Loss: 0.231399... Val Loss: 0.469280
Epoch: 4/4... Step: 1300... Loss: 0.108195... Val Loss: 0.480258
Epoch: 4/4... Step: 1400... Loss: 0.201755... Val Loss: 0.509962
Epoch: 4/4... Step: 1500... Loss: 0.165618... Val Loss: 0.505271
Epoch: 4/4... Step: 1600... Loss: 0.071419... Val Loss: 0.564402
```

Total training time is 76.03 minutes.

4.0.3 Testing

```
[170]: # Get test data loss and accuracy

test_losses = [] # track loss
num_correct = 0

# init hidden state
h = net.init_hidden(batch_size)

net.eval()
# iterate over test data
for inputs, labels in test_loader:

    # Creating new variables for the hidden state, otherwise
    # we'd backprop through the entire training history
    h = tuple([each.detach() for each in h])

    if(train_on_gpu):
        inputs, labels = inputs.cuda(), labels.cuda()

    # get predicted outputs
    output, h = net(inputs, h)

    # calculate loss
    test_loss = criterion(output.squeeze(), labels.float())
    test_losses.append(test_loss.item())

    # convert output probabilities to predicted class (0 or 1)
    pred = torch.round(output.squeeze()) # rounds to the nearest integer

    # compare predictions to true label
    correct_tensor = pred.eq(labels.float().view_as(pred))
```

```

        correct = np.squeeze(correct_tensor.numpy()) if not train_on_gpu else np.
        ↪squeeze(correct_tensor.cpu().numpy())
        num_correct += np.sum(correct)

# -- stats! -- ##
# avg test loss
print("Test loss: {:.3f}".format(np.mean(test_losses)))

# accuracy over all test data
test_acc = num_correct/len(test_loader.dataset)
print("Test accuracy: {:.3f}".format(test_acc))

```

Test loss: 0.578

Test accuracy: 0.788